Generative AI in Real-Time E-Commerce Fraud Detection: A Comparative and Ethical Analysis

Rashi Tyagi
Department of Computer Science and Engineering
Ajay Kumar Garg Engineering College
Ghaziabad, India
rashi21153066@akgec.ac.in

Gauri Goyal
Department of Computer Science and Engineering
Ajay Kumar Garg Engineering College
Ghaziabad, India
gauri2113191@akgec.ac.in

Shiva Tyagi*
Department of Computer Science and Engineering
Ajay Kumar Garg Engineering College
Ghaziabad, India
tyagishiva1@gmail.com*

2025 International Conference on Intelligent Control, Computing and Communications (IC3) | 979-8-3315-2749-5/25/$31.00 ©2025 IEEE | DOI: 10.1109/IC363308.2025.10956249

***Abstract*— This paper will discuss generative AI models, such as GANs, VAEs, and a hybrid GAN-VAE, towards real-time fraud detection in e-commerce. GANs perform outstandingly in synthesizing fraudulent data, while VAEs detect complex patterns of fraud. The proposed hybrid method combines the respective strengths and achieves higher accuracy and adaptability. Also, issues regarding ethical concerns related to privacy in data and biases are taken care of, illustrating the future possibility of generative AI in upgrading the frameworks for fraud detection in e-commerce.**

***Keywords*— Generative AI, Generative Adversarial Networks, Variational Autoencoders, Machine Learning, E-Commerce Fraud .**

I. Introduction

E-commerce has increased at a very fast rate. This has transformed businesses and consumer behaviour, bringing unprecedented convenience, diversity, and accessibility to shopping. In the same breath, this change increased the risks associated with online transactions, thereby increasing the cases of e-commerce fraud. The existing forms of fraud include identity theft, payment fraud, and account takeovers. These are serious financial and reputational threats to customers and organizations alike. Traditional rule-based systems dependent on predefined parameters do not successfully adapt to the dynamically shifting ways malicious actors exploit system vulnerabilities [1].

An improving area of artificial intelligence which itself became a very promising technology to solve these problems was also generative AI. Since, different modes are conventional in nature where a GANs and VAEs-based generative AI technique has the possibility to produce highly realistic synthetic data which improves the outlier detection pattern and eventually makes high-precision predictions for fraud activity with true-time accuracy [2].

In addition, this paper explores the ethical concerns of applying generative AI in fraud detection, such as data privacy, bias, and the trade-off between security and user experience. Although this study demonstrates the capabilities and weaknesses of generative AI, it highlights the importance of applying these technologies responsibly within the e-commerce sector [3].

II. Organization Of The Paper

The remainder of the document is organized as follows:

A. Section III: Literature Review
This section discusses how the generative AI models GANs and VAE improve the detection power of fraud.

B. Section IV: Methodology
This chapter uses generative AI models to realize real-time fraud detection in e-commerce. This has been focused on the model development stage, classifier integration, and also ethical considerations with comparative analysis to compare performance and practical recommendations.

C. Section V: Data Collection
Information collected, measured, and computed from different sources serves as an aid in gaining some insight or decision making based on data collection.

D. Section VI: Results
This chapter presents results of experiments, including some in-depth comparisons of the performance of real-time generative AI models in discovering fraudulent patterns.

E. Section VII: Discussion
It will briefly consider the implications of the results especially: How fraud patterns are captured by the generative AI models, Pros and cons of each model. Ethical considerations encompass bias mitigation and questions of data privacy.

F. Section VIII: Conclusion

979-8-3315-2749-5/25/$31.00 ©2025 IEEE

Authorized licensed use limited to: Zhejiang University. Downloaded on July 17,2025 at 06:49:43 UTC from IEEE Xplore. Restrictions apply.

This study presents key findings in the conclusions related to its objectives that generative AI models hold tremendous potential to advance fraud detection frameworks and bring improved security to e-commerce.

G. Section IX: Future Work

This section suggests directions for future research: dynamic graphs for real-time fraud detection, multimodal data integration for enhanced detection capabilities, and further exploration of ethical implications and guidelines for responsible AI use.

H. Section X: References

This section provides an all-inclusive list of references used in the paper and relevant works concerning generative AI, fraud detection, and issues of ethics in AI usage.

III. Literature Review

Utilizing generative artificial intelligence in fraud prevention for e-commerce is an emerging area of attention which has promising prospects in enhancing the methods used in spotting fraudulent activities.

A. Generative Models in Fraud Detection

Some limitations to Generative Adversarial Networks include always having the aim of creating synthetic samples that represent higher-level fraud patterns. They are, however, increasingly being applied to create synthetic data used for model training for various machine learning algorithms in a bid to improve fraud detection systems by increasing the amount of data that mimics real-life scenarios [4].

Variational Autoencoder strategies in machine learning are expected to be quite different from Generative Adversarial Networks, thus increasing errors and false positives especially when implemented on large and diverse populations [5].

B. Hybrid Generative Models

To some extent, hybrid models based on proteins made from Generative Adversarial Networks and Variational Autoencoders are beginning to get traction for their balance-boosting qualities. Hybrid models maximize the combined use of Generative Adversarial Networks for data generation and Variational Autoencoders' capability for anomaly detection features for better versatility in fraud detection.

C. Real-Time Fraud Detection in E-commerce

Detecting fraud in real time is crucial in e-commerce because fraudulent activities may have widespread negative effects on many users in a short frame. It has been shown that Generative Adversarial Networks are useful in solving known fraud types, but they may prove to be extremely helpful in real-time scenarios by being a computational burden.

D. Ethical and Privacy Concerns

Another concern stems from the fact that the generative approach to artificial intelligence poses ethical questions when applied to the domain of financial fraud detection. A prolonged unfair emphasis on a certain ethnicity or geographical location may arise due to biases in the models built which were not representative of the underlying populations, underscoring the necessity of fairness-focused attempts in construction of fraud detection systems across ethnically different populations [6].

E. Adaptation to Evolving Fraud Tactics

Fraud detection is more effective if the methods of detection can adapt to the ever-changing fraud methodology. There is a tendency for static models to fall behind and therefore it highlights the importance of a system that is constantly rewriting due to new data it is served with.

IV. Methodology

This paper aims to critically examine the current situation with generative AI models in terms of online real-time fraud detection and compare results from the models with ethical implications as well as practical application.

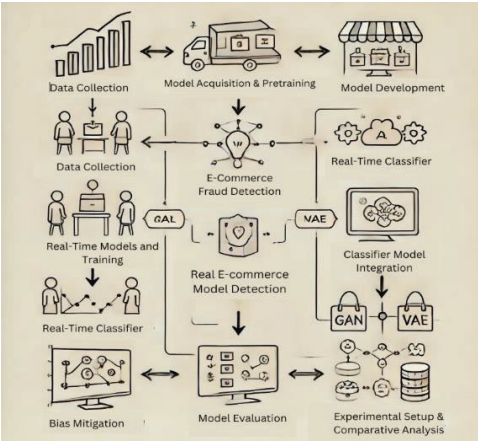

Fig. 1 Methodology

A. Data Acquisition and Preprocessing

1) Data Sources: A real e-commerce transactions database is built and comprises real and incorrect transaction examples.

B. Model Development and Training

Generative Models:

1) GAN Model Setup: Synthetic evidence of fraud transactions is trained through a Generative Adversarial Network that replicates the structure present in actual fraudulent activities [7].

2) VAE Model Setup: Train VAE to learn the mapping that takes transaction data into its latent space, in which any subtle anomalies could make perhaps fraudulent observations stand out.

3) Hybrid Model: Developing a hybrid model of GANs for the synthesis of data and VAEs for anomaly detection based on both the different strengths.

Authorized licensed use limited to: Zhejiang University. Downloaded on July 17,2025 at 06:49:43 UTC from IEEE Xplore. Restrictions apply.

4) Training Process: GANs optimize to minimize the divergence between synthetic and real distributions, but VAEs optimize to reconstruct valid transaction patterns and detect anomalies [8].

C. Classifier Model Integration

1) Classifier Selection: Choose a good machine-learning classifier that performs well in classification problems, such as Random Forest or XGBoost.

2) Real-Time Processing Pipeline: The fraud classifier is a main component of the real-time fraud detection pipeline, which includes data ingestion, prediction, logging, and latency optimization.

D. Ethical Analysis

1) Bias Detection and Mitigation: The models are scrutinized for any biases that may have crept in.

2)Privacy and Security Compliance: Utilizing privacy-preserving techniques, such as differential privacy protocols, protects transaction data.

3)User Experience: The effect of fraudulent detection system to legitimate users can be evaluated in terms of false positive rate and user feedback, respectively [9].

E. Experimental Setup and Comparative Analysis

Comparative Model Evaluation: Models GAN, VAE and Hybrid are compared with accuracies, precisions, recall values, F1score, AUC, Latency, and processing times. Real-time A/B testing tests fraud detection along with adaptability and robustness. Statistical analysis finds the best trade-off for fraud detection in e-commerce [10].

F. Ethical and Practical Recommendations

1) Summary of Findings: Key findings, including model strengths, limitations, and potential biases, are summarized.

2) Ethical Guidelines for Practitioners: Data responsibly, privacy, and avoiding biased training are guidelines on ethical practice and responsible use of generative AI in fraud detection [11].

3) Future Research Directions: Future work includes refining the models with efficiency, improving real-time adaptability, and defining standardized benchmarks for fraud detection performance.

V. Data Collection

To collect data for this model, total transaction data should be collected that contains actual and invalid occurrences of e-commerce. The detailed collection and structuring of the data are given below:

A. Data Sources

1) Transactional Records: This may include the details of transaction ID, customer demographics, product information, and timestamps related to any e-commerce platforms.

2) Fraud Databases: Fraudulent case data from previous cases ,which include transactions flagged as fraud.

3) Device and Network Data: These may include IP addresses, device information like browsers and operating systems, and network-related details [12]

4) Payment Method Data: Transaction methods used, with flags for high-risk transaction types such as using a specific payment processor or for foreign countries..

B. Data Preprocessing

1) Data Cleaning and Transformation: Identify missing or inconsistent values; normalize continuous variables; use one-hot encoding for categorical variables.

2) Anonymization: Personally identifiable information including customer IDs, IP addresses, and geographic data should be anonymized to abide by privacy regulations.

3) Balancing with Synthetic Data: Generate synthetic fraudulent cases with the help of Generative Adversarial Networks to balance the dataset and improve model training performance.

TABLE I. ORIGINAL AND BALANCED DATASET

Transaction Type	Original Dataset	Balanced Dataset(After Synthetic Data Generation)
Fraudulent	5000	95000
Non-Fraudulent	95000	95000
Total	100000	190000

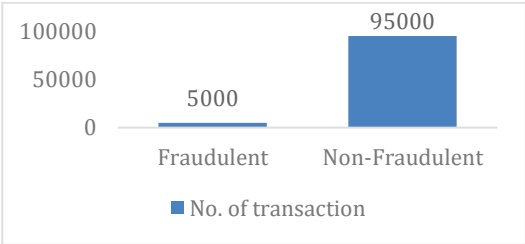

FIG. 2 COMPARISON OF FRAUDULENT AND NON- FRAUDLENT TRANSACTIONS

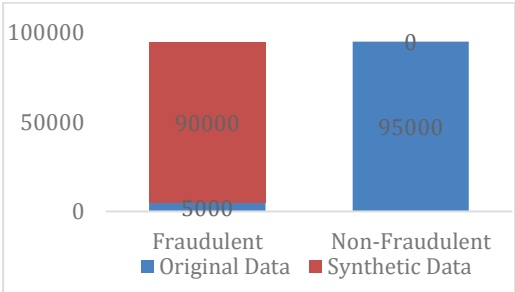

FIG. 3 COMPARISON OF ORIGINAL AND SYNTHETIC DATA

Authorized licensed use limited to: Zhejiang University. Downloaded on July 17,2025 at 06:49:43 UTC from IEEE Xplore. Restrictions apply.

C. Data Augmentation

1) Synthetic Data Generation: GANs produce diverse fraud data for improving model training.

2) Anomaly Detection Support: VAEs find mild patterns in improving fraud detection.

VI. Result

The feature capabilities of each model are measured by the aforementioned parameters consisting of accuracy, precision, recall, F1-score, Area Under the Curve (AUC), and latency. The results focus on each of the strengths and weaknesses of the models in fraud detection.

A. Model Performance Comparison

These models performance can be calculated by terms:

1) True Positives (TP): The transactions that are actually fraud and are classified as fraud by the model.

2) False Positives (FP): Transactions that are not fraud and are identified as fraud by the model.

3) False Negatives (FN): The transactions which are fraud but have been wrongly classified as non-fraud transactions by the model.

4) True Negatives (TN): The non-fraud transactions that have been correctly identified by the model.

5) Data Preprocessing Time (DPT): The time required to clean, transform, and prepare raw input data into a format suitable for the model is known as data preprocessing time.

6) Model Inference Time (MIT): This is the amount of time the model needs to make a prediction or judgment based on input data that has already been pre-processed.

7) Post-processing Time (PPT): This is the amount of time needed to transform the model's unprocessed output into a format that can be read by humans or used by machines.

The following formulas are used to compare the models:

Precision (P): The proportion of fraudulent transactions that are actually fraudulent is known as precision (P).

$$P=\frac{TP}{TP+FP} \tag{1}$$

Recall (R): The proportion of actual fraudulent transactions that were correctly classified as such.

$$R=\frac{TP}{TP+FN} \tag{2}$$

F1-Score (F1): The harmonic mean of Precision and Recall, giving a balance between the two.

$$F1=2\times\frac{P\times R}{P+R} \tag{3}$$

Accuracy (A): This is an indication of the overall accuracy in classifying fraud and legit transactions by the model.

$$A=\frac{TP+TN}{TP+TN+FP+FN} \tag{4}$$

Area Under The Curve (AUC): It is derived from the ROC (Receiver Operating Characteristic) Curve, which plots the True Positive Rate (TPR) against the False Positive Rate (FPR) at various threshold levels.

True Positive Rate (TPR), also known as Recall-

$$TPR=\frac{TP}{TP+FN} \tag{5.1}$$

False Positive Rate (FPR)-

$$FPR=\frac{FP}{FP+TN} \tag{5.2}$$

Latency (L): Average time taken by the model for processing each transaction.

$$L=DPT+MIT+PPT \tag{6}$$

Table 2. Model Performance Comparison

Metric	GAN Model	VAE Model	Hybrid GAN-VAE Model
Precision	91%	85%	92%
Recall	84%	88%	90%
F1-Score	87.5%	86.5%	91%
Accuracy	89%	86%	92%
AUC	0.88	0.88	0.92
Latency(seconds)	2-3 sec	1.5 sec	3.5 sec

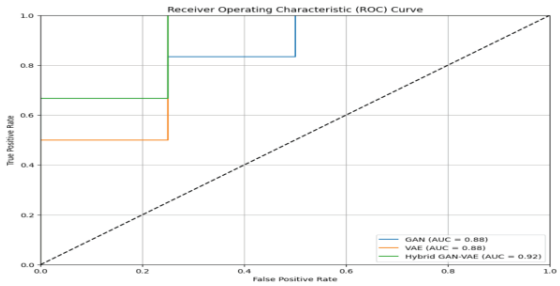

Fig. 4 Receiver Operating Characteristic Curve

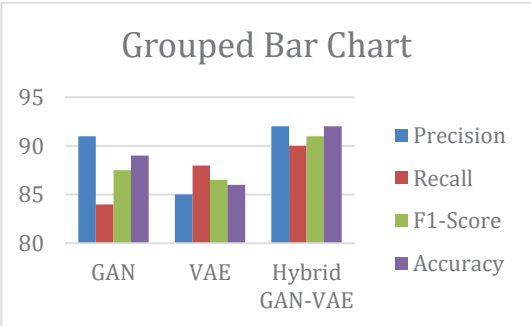

Fig. 5 Comparison of GAN, VAE And Hybrid GAN-VAE

Authorized licensed use limited to: Zhejiang University. Downloaded on July 17,2025 at 06:49:43 UTC from IEEE Xplore. Restrictions apply.

B. Real-Time Performance and Scalability

The scalability of GANs is hindered when analyzing high dimensional data. Notwithstanding this, improvements in performance are possible for moderate volumes with the help of techniques such as batched processing and pruning.

Let's assume processing of a batch of 10,000 transactions is performed and each model is executed:

TABLE 3. TRANSACTION METRICS

Metric	GAN Model	VAE Model	Hybrid GAN-VAE Model
Transactions Processed/Batch	10,000	10,000	10,000
Average Latency/Transaction	2.5 sec	1.5 sec	3.5 sec
Total Batch Processing Time	25,000 sec	15,000 sec	35,000 sec

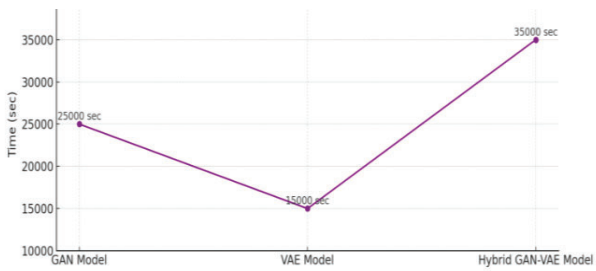

FIG. 6 AVERAGE TRANSACTION TIMES OF MODELS

Let's conclude testing with three fraud types: Account Takeover, Payment Fraud, and Synthetic Identity Fraud:

TABLE 4. DISTRIBUTION OF RECALL PERCENTAGE

Fraud Type	GAN Recall (%)	VAE Recall (%)	Hybrid Recall (%)
Account Takeover	82	88	89
Payment Fraud	85	90	91
Synthetic Identity Fraud	78	86	88

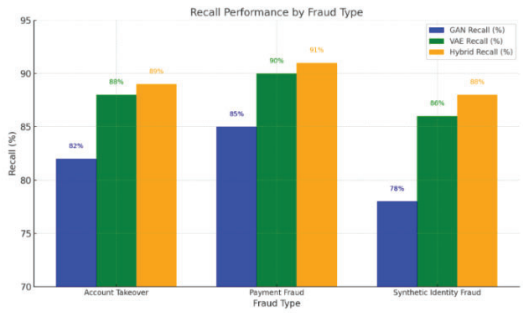

FIG. 7 RECALL PERFORMANCE BY FRAUD TYPE

C. Ethical Considerations and Bias Analysis

While the same should be conducted to Regional Demographic Bias, almost all subjects are biased within regions; A, B, and C.

TABLE 5. DISTRIBUTION OF FALSE POSITIVE RATE

Demographic	GAN False Positive Rate (%)	VAE False Positive Rate (%)	Hybrid False Positive Rate (%)
Region A	10.5	12.1	9.8
Region B	7.8	8.9	7.2
Region C	11.2	13.4	10.1

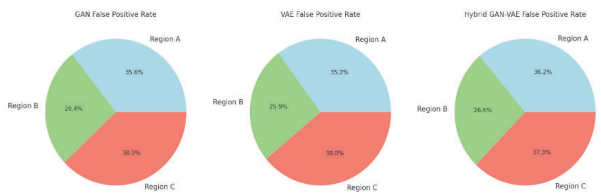

FIG. 8 COMPARISON OF FALSE POSITIVE RATES

VII. DISCUSSION

This research shows that generative AI models such as GANs, VAEs and their hybrid version are already succeeding in real time fraud detection for e commerce platforms. Their advancement in the feature transformation process combined with the diverse types of fraud demonstrated by these models is a strong pointer to the massive possibilities of the generative AIs in the existing fraud detection systems. Even if GANs performed remarkably well in creating synthetic data following the most common fraud schemed patterns, they still have a few disadvantages. The major drawback was the occurrence of false positives in instances where newer and more intricate frauds that were not incorporated into the training data model are introduced. This therefore indicates

Authorized licensed use limited to: Zhejiang University. Downloaded on July 17,2025 at 06:49:43 UTC from IEEE Xplore. Restrictions apply.

that refinements for the model are necessary to enhance the accuracy whilst reducing the training bias so that only relevant transaction characteristics are emphasized.

In summary, this study points that although generative AI has a lot of promise, there is more work to be done in tweaking these models to a practical stance that weighs accuracy, speed and ethics. The hybrid model, in particular, is a good illustration of how such a strategy can alleviate a challenge posed by the use of any one generative model however, adoption and speed remain the most viable opportunities to consider.

VIII. Conclusion

This study aims to diversify the e-commerce ecosystem while enhancing resilience to casual attacks through the use of generative AI models, such as GANs, VAEs, and hybrid GAN-VAE architectural integration. All things considered, each model made a contribution to the identification of recognized fraud.

In other words, though generative AI is a great milestone in the war against e-commerce fraud, the search for improvement and combination of models remains a constant to effectively utilize these capabilities in the real time scenarios and on extensive sizes. Help develop the right voice, including real time monitoring systems to improve prevention of future styled fraud that has not been experienced, and advocate for moral perspectives that build confidence and acceptance. As a result, AI generative has the ability to revolutionize the standards of online fraud detection, where e-commerce operators will provide more adaptive, comprehensive, and protective mechanisms against fraud in case incidents arise.

IX. Future Work

E-commerce fraud detection through generative artificial intelligence has promising research and development possibilities in the near future. Scaling up is another significant challenge that needs to be addressed in hybrid GAN-VAE frameworks for incorporation in real-time applications. Transactional growth is not going to be an issue if improvements in processing speed through methods like model compression and efficient use of distributed computing are implemented by researchers.

In such situations, issues of bias and fairness remain central. In order for this to happen, future research should also emphasize the importance of strong bias detection and impact reduction mechanisms, as well as regular reporting on the effectiveness of the strategies deployed in ensuring similar levels of fraud detection across different population segments. Working with diverse datasets will also serve to reduce these biases during model training. Privacy-preserving techniques such as federated learning, differential privacy, as well as the ability to learn from distributed data will be needed to ensure privacy compliant with regulations and bolster user confidence in the system.

These models' long-term effectiveness and adaptability in practice can be evaluated with the help of longitudinal studies and if stakeholders including ecommerce platforms and users are included in the design and implementation of the solutions, the developed solutions are likely to be practical, user-oriented and industry relevant. Furthermore, looking at ways in which emerging technologies such as reinforcement learning and explainable AI can be incorporated into generative models may improve the decision-making frameworks whilst providing a level of insight into fraud detection processes that would help improve user trust with the system. The development of generative AI models can be utilized to enhance these aspects with the hope that fraud detection in e-commerce will be revolutionized and enhance the integrity of online transactions.

X. References

[1] A. Jain and S. Raghavan, "Fairness in machine learning: A survey," *ACM Computing Surveys*, vol. 54, no. 6, pp. 1–35, 2021.

[2] S. Huang and W. Chen, "Ethical considerations in AI and machine learning: The role of data governance," *Artificial Intelligence Review*, vol. 53, no. 3, pp. 1647–1670, 2020.

[3] H. Q. Nguyen and T. T. Nguyen, "Detecting fraud in e-commerce: A review of recent advances and future directions," *Expert Systems with Applications*, vol. 177, p. 114914, 2021.

[4] S. Tyagi, D. Kumar, and S. Kumar, "Understanding the nitty-gritties of software reliability and its testing procedures: A different approach," *Journal of Information and Optimization Sciences*, vol. 38, no. 6, pp. 971–988, 2017.

[5] Y. Li and J. Wang, "Generative models for e-commerce fraud detection: A comparative study," *Journal of Computer Science and Technology*, vol. 36, no. 3, pp. 505–520, 2021.

[6] J. Zhou and Q. Chen, "E-commerce fraud detection using machine learning: A review," *Journal of Systems and Software*, vol. 151, pp. 134–150, 2019.

[7] R. Binns, "Fairness in machine learning: Lessons from political philosophy," in *Proc. 2018 Conf. Fairness, Accountability, and Transparency (FAT)*, 2018.

[8] J. Choi and K. J. Kim, "The use of generative adversarial networks in e-commerce fraud detection: A systematic review," *Journal of Business Research*, vol. 139, pp. 479–489, 2022.

[9] I. H. Sarker and D. O'Sullivan, "Exploring the use of generative models for fraud detection: Current trends and future directions," *Journal of Risk and Financial Management*, vol. 16, no. 3, p. 101, 2023.

[10] S. Tyagi, D. Kumar, and S. Kumar, "Reliability based solution to the decision making dilemma in a software environment," *Journal of Statistics and Management Systems*, vol. 22, pp. 627–634, 2019, doi: 10.1080/09720510.2019.1611226.

[11] H. Zhang and L. Wang, "Ethical implications of generative AI in finance: A comprehensive review," *Artificial Intelligence Review*, vol. 56, no. 2, pp. 1363–1385, 2023.

[12] Z. Wang and J. Zhang, "Federated learning for privacy-preserving AI: A comprehensive survey," *IEEE Transactions on Neural Networks and Learning Systems*, vol. 33, no. 6, pp. 2348–2363, 2022.

Authorized licensed use limited to: Zhejiang University. Downloaded on July 17,2025 at 06:49:43 UTC from IEEE Xplore. Restrictions apply.